

Forecasting NBA 3 Point Shot Percentage

```
# Load packages:
```

```
library(MASS)
library(forecast)
```

```
## Registered S3 method overwritten by 'quantmod':
##   method      from
##   as.zoo.data.frame zoo
```

```
library(tseries)
library(astsa)
```

```
##
## Attaching package: 'astsa'
```

```
## The following object is masked from 'package:forecast':
##
##   gas
```

```
library(dse)
```

```
## Loading required package: tfplot
```

```
## Loading required package: tframe
```

```
##
## Attaching package: 'dse'
```

```
## The following objects are masked from 'package:forecast':
##
##   forecast, is.forecast
```

```
## The following objects are masked from 'package:stats':
##
##   acf, simulate
```

```
library(knitr)
library(gridExtra)
```

```
##
## Attaching package: 'gridExtra'
```

```
## The following object is masked from 'package:dse':  
##  
##      combine
```

```
library(grid)  
library(tidyverse)
```

```
## -- Attaching packages -----
```

```
## v ggplot2 3.3.1      v purrr  0.3.4  
## v tibble  3.0.1      v dplyr  1.0.0  
## v tidyr   1.1.0      v stringr 1.4.0  
## v readr   1.3.1      v forcats 0.5.0
```

```
## -- Conflicts -----
```

```
## x dplyr::combine() masks gridExtra::combine(), dse::combine()  
## x dplyr::filter()  masks stats::filter()  
## x dplyr::lag()     masks stats::lag()  
## x dplyr::select()  masks MASS::select()  
## x purrr::splice()  masks tframe::splice()
```

```
library(MAPA)
```

```
## Loading required package: parallel
```

```
## Loading required package: RColorBrewer
```

```
## Loading required package: smooth
```

```
## Loading required package: greybox
```

```
## Package "greybox", v0.6.0 loaded.
```

```
##
```

```
## Attaching package: 'greybox'
```

```
## The following object is masked from 'package:tidyr':
```

```
##
```

```
##      spread
```

```
## The following objects are masked from 'package:dse':
```

```
##
```

```
##      forecast, polyprod
```

```
## This is package "smooth", v2.6.0
```

```
##
```

```
## Attaching package: 'smooth'
```

```
## The following object is masked from 'package:dse':
```

```
##
```

```
##      forecast
```

```

lbjmv<-read.csv(file = 'sportsref_download.csv',
               header = TRUE, sep = ',', skip = 1)

#Remove last few entries for training.
trainset <- lbjmv %>%
  select(Season, X3PA, FGA) %>%
  mutate(percentageof3=X3PA/FGA) %>%
  filter(Season!='2019-20',Season!='2018-19',Season!='2017-18') %>%
  filter(!is.na(X3PA)) %>%
  arrange(-row_number())
#length(trainset)

#Full data. Subtract 2019-2020 season, since it's not over yet (at time
#of project!)
testset <- lbjmv %>%
  select(Season, X3PA, FGA) %>%
  mutate(percentageof3=X3PA/FGA) %>%
  filter(Season!='2019-20') %>%
  filter(!is.na(X3PA)) %>%
  arrange(-row_number())


trainset<-trainset[,4]
#trainset

testset<-testset[,4]
#testset

ts.plot(testset, ylab = "Percentage of Shots Taken that are 3",
        xlab = "Years since 1979")
title(expression(Percentage~of~3~Point~Shots~Taken~From~1979~--~Present))


```

Percentage of 3 Point Shots Taken From 1979 – Present




#Data does not appear to display seasonality.


```
trainset.ts <- ts(trainset, frequency = 1)
testset.ts <- ts(testset, frequency = 1)
#Initial look at ACF and PACF plots.
acf(trainset.ts, lag.max = 20, main = "" )
```



```
pacf(trainset.ts, lag.max=20, main = "")
```



```
#Perhaps variance needs stabilizing?  
bctrainset <- boxcox(trainset.ts~as.numeric(1:length(trainset)))
```




```
lambda1 <- bctrainset$x[which.max(bctrainset$y)]
lambda1
```

```
## [1] 0.989899
```

```
#Lambda is 0.989899
trainset.tr <- trainset.ts^(0.989899)
ts.plot(trainset.tr, ylab = "Three Point FGA BoxCox",
        xlab = "Years Since 1979",
        main = "Box Cox Transformed Data")
```

Box Cox Transformed Data



```
var(trainset.ts)
```

```
## [1] 0.007339486
```

```
#0.007339486
```

```
var(trainset.tr)
```

```
## [1] 0.007496046
```

```
#0.007496046
```

```
#Slight increase in variance, but we will allow this for now.
```

```
#There appears to be no seasonality, so we will difference to remove trend.
```

```
trainsetdiff1 <- diff(trainset.tr, lag =1) #Difference once and observe.
```

```
var(trainsetdiff1)
```

```
## [1] 0.0002683378
```

```
#0.0002683378
```

```
#ts.plot(trainsetdiff1, ylab = "Differenced At Lag 1")
```

```
#We want it to resemble white noise.
```

```
#abline(lm(trainsetdiff1~as.numeric(1:length(trainsetdiff1))),
```




```
#      col = "red")

trainsetdiff1diff1 <- diff(trainsetdiff1, lag = 1)
var(trainsetdiff1diff1)

## [1] 0.0004814831

#0.0004814831
#ts.plot(trainsetdiff1diff1, ylab = "Differenced At Lag 1")
#abline(lm(trainsetdiff1diff1~as.numeric(1:length(trainsetdiff1diff1))),
#      col = "red")

par(mfrow=c(1,2))
ts.plot(trainsetdiff1, ylab = "Differenced At Lag 1")
abline(lm(trainsetdiff1~as.numeric(1:length(trainsetdiff1))),
      col = "red")
ts.plot(trainsetdiff1diff1, ylab = "Differenced At Lag 1, Lag 1")
abline(lm(trainsetdiff1diff1~as.numeric(1:length(trainsetdiff1diff1))),
      col = "red")
```



```
#Differencing twice may work!
```

```
#Perhaps more differencing may be useful?  
#In general, it's best to not overdifference...
```


#but this may provide useful insight!

```
trainsetdiff1diff1diff1 <- diff(trainsetdiff1diff1, lag =1)
var(trainsetdiff1diff1diff1)
```

```
## [1] 0.001511063
```

#0.001511063 Too high?

```
ts.plot(trainsetdiff1diff1diff1,
        ylab = "Differenced At Lag 1, Lag 1, Lag 1")
abline(lm(trainsetdiff1diff1diff1~
          as.numeric(1:length(trainsetdiff1diff1diff1))),
       col = "red")
```




```
trainsetdiff1diff1diff1diff1 <- diff(trainsetdiff1diff1diff1, lag =1)
var(trainsetdiff1diff1diff1diff1)
```

```
## [1] 0.005547701
```

#0.005547701 Too high?


```
ts.plot(trainsetdiff1diff1diff1diff1,
        ylab = "Differenced At Lag 1, Lag 1, Lag 1, Lag 1")
abline(lm(trainsetdiff1diff1diff1diff1~
          as.numeric(1:length(trainsetdiff1diff1diff1diff1))),
       col = "red")
```



```
trainsetdiff1diff1diff1diff1diff1 <- diff(trainsetdiff1diff1diff1diff1, lag =1)
var(trainsetdiff1diff1diff1diff1diff1)
```



```
## [1] 0.02145416
```

```
#0.02145416 Too high!
ts.plot(trainsetdiff1diff1diff1diff1diff1,
        ylab = "Differenced At Lag 1, Lag 1, Lag 1, Lag 1, Lag 1")
abline(lm(trainsetdiff1diff1diff1diff1diff1~
          as.numeric(1:length(trainsetdiff1diff1diff1diff1diff1))),
       col = "red")
```




```
#The data may be overdifferenced, but it clearly resembles white noise!  
#More analysis is needed.
```

```
# Differenced at lag 1 PACF AND ACF  
par(mfrow=c(1,2))  
acf(trainsetdiff1, lag.max = 24, main = "")  
pacf(trainsetdiff1, lag.max = 24, main = "")
```



```
# DIFFERENCED at lag 1 and lag 1 and lag 1 PACF AND ACF
par(mfrow=c(1,2))
acf(trainsetdiff1diff1diff1, lag.max=24, main = "")
pacf(trainsetdiff1diff1diff1, lag.max=24, main = "")
```



```
# S = 0, no seasonality.
# D = 0, d = 3, differenced 3x to remove trend
# Examine at lag = 1,2,3,4,5,...
# ACF plot tails off (q = 0)
# and PACF cuts off after lag 4 (p = 4) ----> ARIMA(4,3,0)?
# Or, perhaps ACF cuts off at lag 4 (q = 4)
# and PACF tails off (p = 0) -> ARIMA(0,3,4)?
# More analysis is needed.
# We will consider multiple models and choose the best one.
```

```
#We want the lowest possible AICc values.
```

```
mod <- Arima(trainset.tr, order = c(4,3,0)) #AICc ARIMA(4,3,0) -160.6
mod
```

```
## Series: trainset.tr
## ARIMA(4,3,0)
##
## Coefficients:
##      ar1      ar2      ar3      ar4
##    -1.0734 -0.6138 -0.692 -0.4005
## s.e.   0.1508   0.2064   0.199   0.1456
##
## sigma^2 estimated as 0.0004503: log likelihood=86.32
## AIC=-162.64  AICc=-160.57  BIC=-154.86
```

```
mod2 <- Arima(trainset.tr, order = c(0,3,4)) #AICc ARIMA(0,3,4) -169.5
mod2
```

```
## Series: trainset.tr
## ARIMA(0,3,4)
##
## Coefficients:
##          ma1      ma2      ma3      ma4
##      -1.4818  0.6229 -0.5113  0.4982
## s.e.   0.2273  0.2872  0.2288  0.1685
##
## sigma^2 estimated as 0.0003048: log likelihood=90.77
## AIC=-171.54  AICc=-169.47  BIC=-163.76
```

#Slight adjust to consider more models.

```
mod3 <- Arima(trainset.tr, order = c(4,2,0)) #AICc ARIMA(4,2,0) -182.1
mod3
```

```
## Series: trainset.tr
## ARIMA(4,2,0)
##
## Coefficients:
##          ar1      ar2      ar3      ar4
##      -0.5714 -0.1501 -0.4958 -0.2705
## s.e.   0.1602  0.1664  0.1600  0.1545
##
## sigma^2 estimated as 0.0002915: log likelihood=97.03
## AIC=-184.05  AICc=-182.05  BIC=-176.14
```

```
mod4 <- Arima(trainset.tr, order = c(0,2,4)) #AICc ARIMA(0,2,4) -186.2
mod4
```

```
## Series: trainset.tr
## ARIMA(0,2,4)
##
## Coefficients:
##          ma1      ma2      ma3      ma4
##      -0.7965  0.0086 -0.5853  0.3733
## s.e.   0.2193  0.2015  0.1677  0.2042
##
## sigma^2 estimated as 0.0002378: log likelihood=99.11
## AIC=-188.21  AICc=-186.21  BIC=-180.3
```

*#We have multiple models to consider. But the auto.arima() function
#may give us an even better recommendation.*

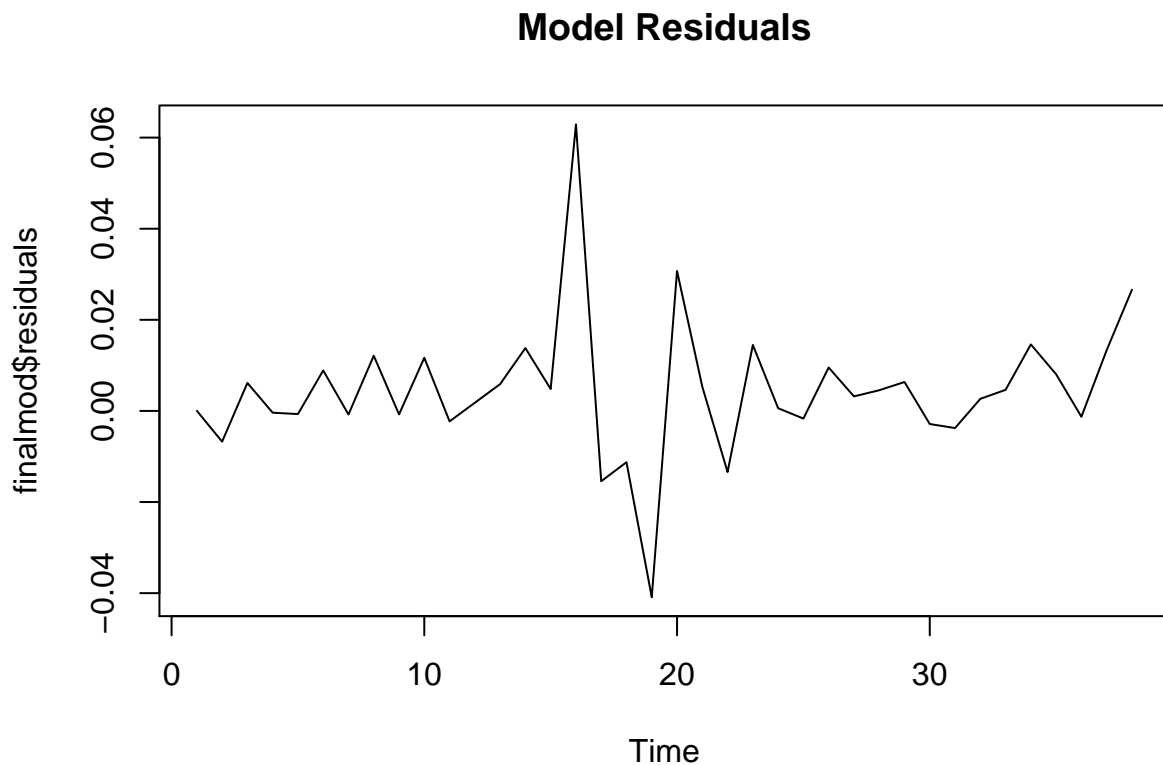
```
modauto<-auto.arima(trainset.tr,seasonal = FALSE,
                    stepwise = FALSE,
                    approximation = FALSE,
                    allowdrift = FALSE)
modauto #AICc ARIMA(0,1,2) -193.8 Lowest one yet!
```

```
## Series: trainset.tr
## ARIMA(0,1,2)
##
## Coefficients:
##      ma1      ma2
##    0.4066 0.4651
## s.e. 0.1826 0.1579
##
## sigma^2 estimated as 0.0002698: log likelihood=100.27
## AIC=-194.55 AICc=-193.82 BIC=-189.72
```

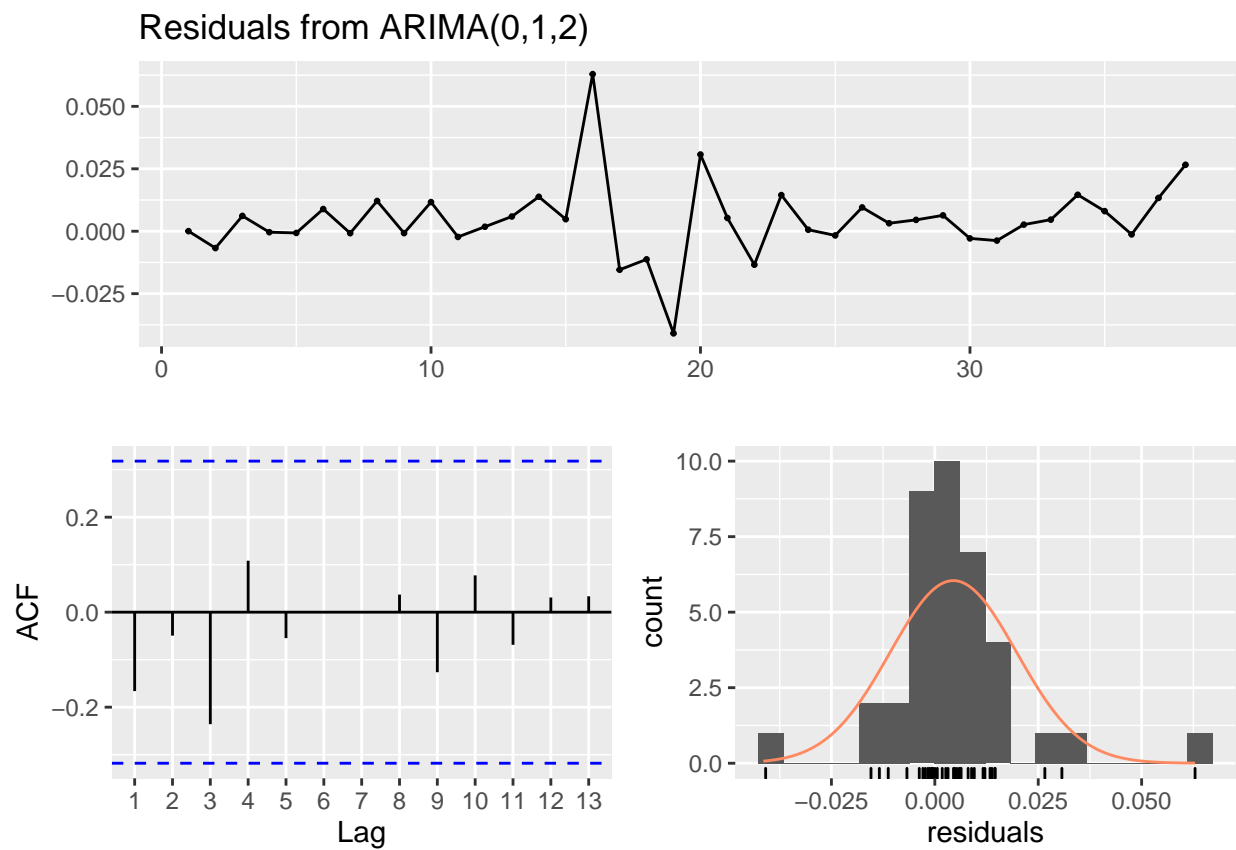
```
finalmod <- Arima(trainset.tr, order = c(0,1,2))
finalmod
```

```
## Series: trainset.tr
## ARIMA(0,1,2)
##
## Coefficients:
##      ma1      ma2
##    0.4066 0.4651
## s.e. 0.1826 0.1579
##
## sigma^2 estimated as 0.0002698: log likelihood=100.27
## AIC=-194.55 AICc=-193.82 BIC=-189.72
```

```
ts.plot(finalmod$residuals, main = "Model Residuals")
```





```
#Resembles white noise.
checkresiduals(finalmod)
```




```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(0,1,2)
## Q* = 4.3833, df = 6, p-value = 0.625
##
## Model df: 2.    Total lags used: 8
```

```
#Residuals are normally distributed.
autoplot(finalmod)
```



```
#Roots lie within unit circle. Implies stationarity and invertibility.  
qqnorm(finalmod$residuals, main = "Normal QQ Plot for Model")  
qqline(finalmod$residuals)
```

Normal QQ Plot for Model



#Residuals are normally distributed.

```
modelpredictions <- predict(finalmod, n.ahead=2)

upperbound <- modelpredictions$pred + 2*modelpredictions$se
lowerbound <- modelpredictions$pred - 2*modelpredictions$se

ts.plot(trainset,
        xlim=c(1, length(trainset)+5),
        ylim=c(0,0.5),
        main = "Forecasting on Data",
        ylab= "% of 3PT FGA")

lines(upperbound, col="blue", lty = "dashed")
lines(lowerbound, col="blue", lty = "dashed")

points((length(trainset)+1):(length(trainset)+2),
       modelpredictions$pred, col = "red")

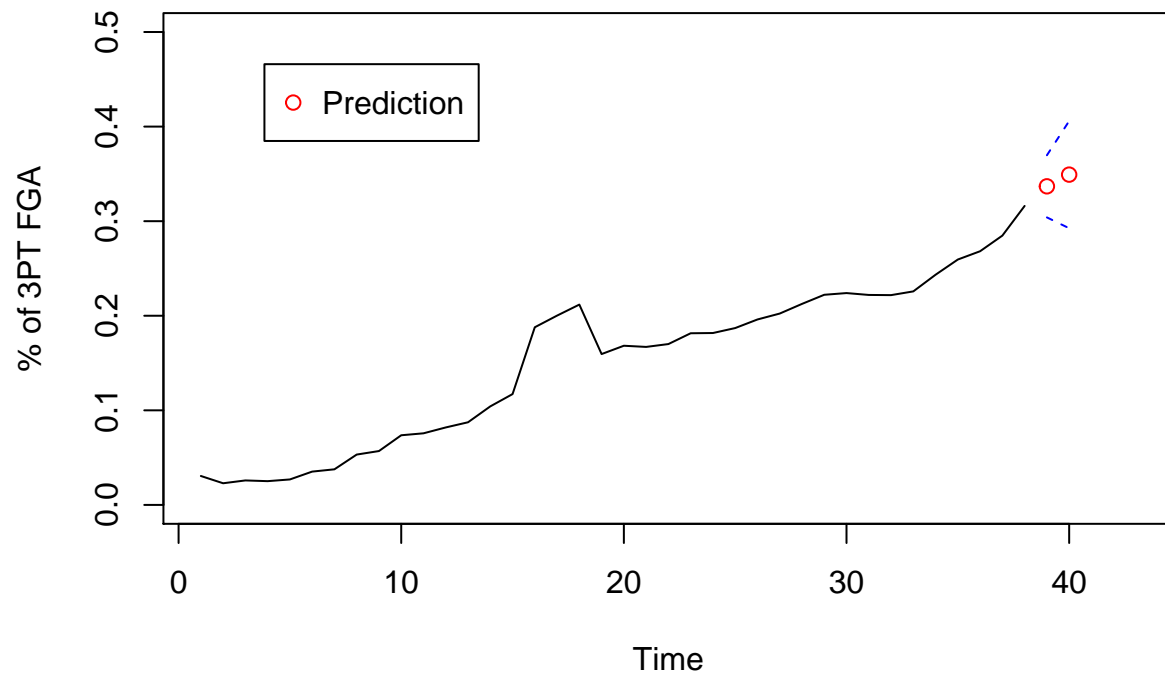
# Legend:
legend("topleft",
      legend = c("Prediction"),
      col = c("red"),
      pch = 1,
      bty = "o",
      pt.cex = 1,
```

```

cex = 1,
text.col = "black",
horiz = F ,
inset = c(0.1, 0.1))

```

Forecasting on Data



```

#Zoom in! Compare it to test set.

ts.plot(testset,
        xlim = c(length(testset)-5, length(testset)+3),
        ylim = c(0.15,0.45),
        main = "Observed vs Forecasted Values",
        ylab = "3PTFGA%")

# Points for observed data
points((length(testset)+1):(length(testset)+2),
       testset.ts[39:40], col = "blue")

# Points for forecasted data
points((length(testset)+1):(length(testset)+2),
       modelpredictions$pred, col = "red")

lines((length(testset)+1): (length(testset)+2),
      upperbound, lty=2, col = "blue")

lines((length(testset)+1): (length(testset)+2),

```


```

lowerbound, lty=2, col = "blue")

# Add a legend
legend("bottomright",
      legend = c("Prediction", "Observed"),
      col = c("red",
              "blue"),
      pch = 1,
      bty = "o",
      pt.cex = 1,
      cex = 1,
      text.col = "black",
      horiz = F ,
      inset = c(0.1, 0.1))

```

Observed vs Forecasted Values



```

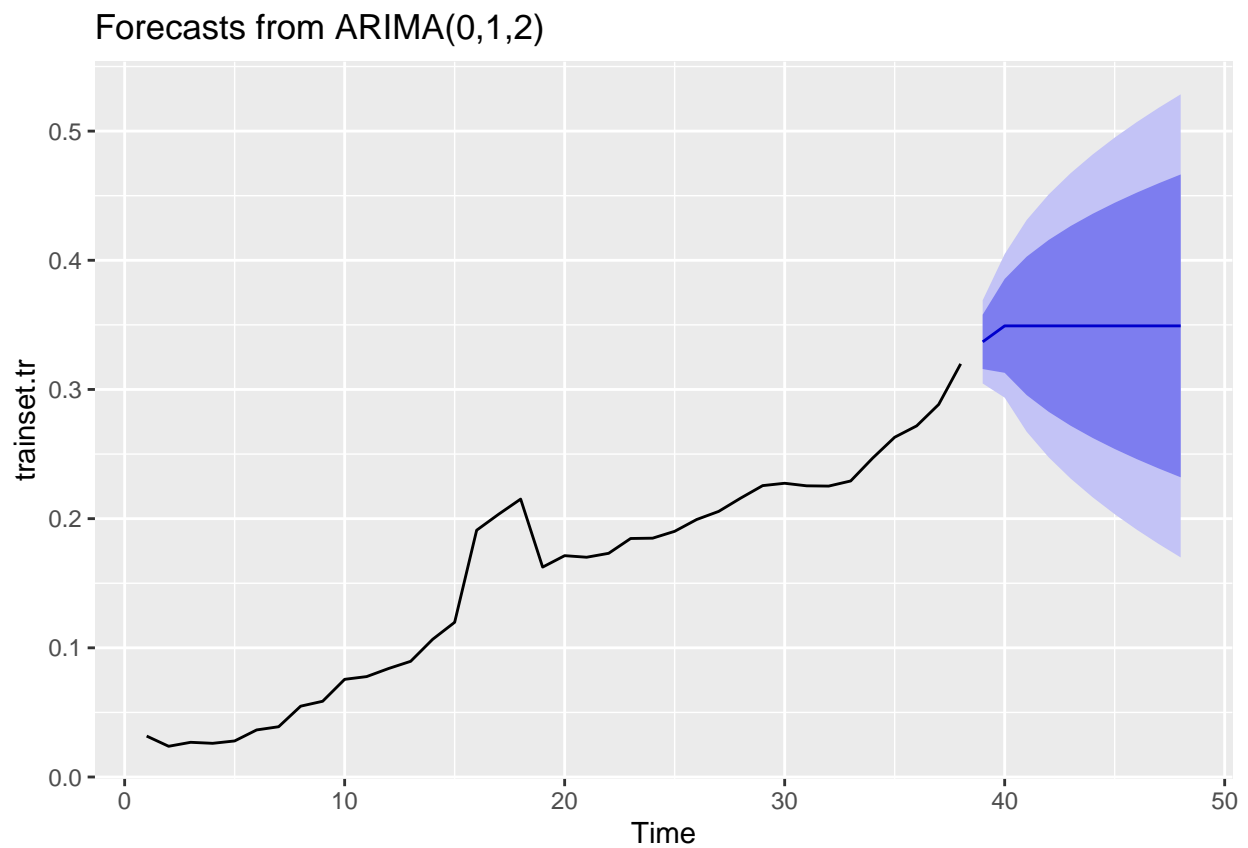
predict(finalmod, n.ahead = 5) #Our ARIMA(0,1,2)

## $pred
## Time Series:
## Start = 39
## End = 43
## Frequency = 1
## [1] 0.3368486 0.3492225 0.3492225 0.3492225 0.3492225
##
## $se

```

```
## Time Series:
## Start = 39
## End = 43
## Frequency = 1
## [1] 0.01642527 0.02834745 0.04181777 0.05190255 0.06032432
```

```
library(MAPA)
#install.packages('smooth')
library(smooth)
modelforecasts<-forecast(finalmod)
library(forecast)
autoplot(modelforecasts)
```



#Our model suggests the data converges to a point! Namely, 35% 3PT FGA.